Scaling Enumerative Program Synthesis via Divide and Conquer

Abstract. Given a semantic constraint specified by a logical formula, and syntactic constraints specified by a context-free grammar, the Syntax-Guided Synthesis (SyGuS) problem is to find an expression that satisfies both the syntactic and semantic constraints. An enumerative approach to solve this problem is to systematically generates all expressions from the syntactic space with some pruning, and has proved to be surprisingly competitive in the newly started competition of SyGuS solvers. It performs well on small to medium sized benchmarks, produces succinct expressions, and has the ability to generalize from input-output examples. However, its performance degrades drastically with the size of the smallest solution. To overcome this limitation, in this paper we propose an alternative approach to solve SyGuS instances.

The key idea is to employ a divide-and-conquer approach by separately enumerating (a) smaller expressions that are correct on subsets of inputs, and (b) predicates on inputs that distinguish these subsets. These smaller expressions and predicates are then combined using decision trees to obtain an expression that is correct on all inputs. We view the problem of combining expressions and predicates as a multi-label decision tree learning problem. We propose a novel technique of associating a probability distribution over the set of labels that a sample can be labeled with. This enables us to use standard information-gain based heuristics to learn a compact decision tree.

We report a prototype implementation and evaluate it on the benchmarks from the SyGuS 2015 competition. Our tool is able to match the running times and the succinctness of solutions of both standard enumerative solver as well as the latest white-box solvers in most cases. Further, our solver is able to solve a large number of instances from the ICFP class of benchmarks, which were previously unsolved by all existing solvers.

1   Introduction

The field of program synthesis relates to automated techniques that attempt to automatically generate programs from requirements that a programmer writes. It has been applied to various domains such as program completion [23], program optimization, and automatic generation of string manipulation programs from input-output examples [8], among other applications. Recently, Syntax-Guided Synthesis (SyGuS) has been proposed as a back-end exchange format and enabling technology for program synthesis [2]. The aim is to allow experts from different domains to model their synthesis problems as SyGuS instances, and leverage general purpose SyGuS solvers.

In the SyGuS approach, a synthesis task is specified using restrictions on both the form (syntax) and function (semantics) of the program to be synthesized: (a) The syntactic restrictions are given in terms of a context-free grammar from

which a solution program may be drawn. (b) The semantic restrictions are encoded into a specification as an SMT formula. Most SyGuS solvers operate in two cooperating phases: a *learning phase* in which a candidate program is proposed, and a *verification phase* in which the proposal is checked against the specification. SyGuS solvers can be broadly categorized into two kinds: (a) black-box solvers, where the learning phase does not deal with the specification directly, but learns from constraints on how a potential solution should behave on sample inputs points [2, 20, 24]; and (b) white-box solvers, which attempt learn directly from the specification, generally using constraint solving techniques [3, 19].

The enumerative solver [2] placed first and second in the SyGuS competition 2014 and 2015, respectively. It maintains a set of concrete input points, and in each iteration attempts to produce an expression that is correct on these concrete inputs. It does so by enumerating expressions from the grammar and checking if they are correct on the input points, while pruning away expressions that behave equivalently to already generated expressions. If an expression that is correct on the input points is found, it is verified against the full specification. If it is incorrect, a counter-example point is found and added to the set of input points.

While the enumerative strategy works well when the desired solutions have small to medium sizes, it does not scale well with the solution size. Figure 2 shows that the time taken to explore all potential solutions grows exponentially with solution size. To overcome this scalability issue, we introduce a divide-and-conquer enumerative algorithm.

The divide-and-conquer enumerative approach is based on this insight: while the full solution expression to the synthesis problem may be large, the important individual parts are small. The individual parts we refer to here are: (a) *terms* which serve as the return value for the solution, and (b) *predicates* which serve as the conditionals that choose which term is the actual return value for a given input. For example, in the expression if $x \leq y$ then y else x, the terms are x and y, and the predicate is $x \leq y$. In this example, although the full expression has size 6, the individual terms have size 1 each, and the predicate size 3. Hence, the divide-and-conquer enumerative approach only enumerates terms and predicates separately and attempts to combine them into a conditional expression.

For combining the different part of the solution in a full conditional expression, we use the technique of learning decision trees [4, 18]. The input points maintained by the enumerative algorithm serve as the samples, the predicates enumerated serve as the attributes, and the terms serve as the labels. A term t is a valid label for a point pt if t is correct for pt. We use a simple multi-label decision tree learning algorithm returns a sound decision tree that classifies the samples soundly, i.e., for each point, following the edges corresponding to the attribute values (i.e., predicates) lead to a branch.

To enhance the quality of the solutions obtained, we extend the basic divide-and-conquer algorithm to be an *anytime* algorithm, *i.e.*, the algorithm does not stop when the first solution is found, and instead continues enumerating terms and predicates in an attempt to produce more compact solutions. Decomposing the verification queries into *branch-level queries* helps in faster convergence.

Round no.	Enumerated expressions	Candidate Expression	Point added
1	0	0	$\{x \mapsto 1, y \mapsto 0\}$
2	0, 1	1	$\{x \mapsto 0, y \mapsto 2\}$
3	0, 1, x, y, ..., $x + y$,	$x + y$	$\{x \mapsto 1, y \mapsto 2\}$
$\vdots$	$\vdots$	$\vdots$	$\vdots$
n	0, ..., `if` $x \le y$ `then` y `else` x	`if` $x \le y$ `then` y `else` x	

Table 1: Example run of the basic enumerative algorithm

Evaluation. We implemented the proposed algorithm in a prototype tool and evaluated it on benchmarks from the SyGuS competition 2015, both in the linear integer arithmetic and the bit-vector domain. The tool was able to perform on par or better than existing solvers in the bit-vector domain, and produced compact solutions. The tool was able to solve 47 out of the 50 problems in the ICFP suite of benchmarks. To the best of our knowledge, no other solver has been able to solve the ICFP benchmarks satisfactorily. Further, we also observed that the anytime extension of the algorithm was able to produce more compact solutions in 22 of the 47 solved benchmarks. We were able to obtain upto 90% reduction in solution size in some cases. In the linear integer arithmetic domain, the tool performed much better than existing black-box solvers, scaling to larger instances. However, for larger benchmarks, the performance was worse, and the solution size was slightly larger than those produced by white-box solvers due to the shortcomings of the decision tree learning algorithm.

2 Illustrative Example

We explain the salient points of our approach on an example. Consider a synthesis task to generate an expression e such that: (a) e is generated by the grammar from Figure 1. (b) e when substituted for f, in the specification Φ,

```
S ::= T | if (C) then T else T
T ::= 0 | 1 | x | y | T + T
C ::= T ≤ T | C ∧ C | ¬ C
```
Fig. 1: Grammar for linear integer expressions

renders it true, where $\Phi \equiv \forall x, y : f(x, y) \ge x \wedge f(x, y) \ge y \wedge (f(x, y) = x \vee f(x, y) = y)$. Note that the specification constrains $f(x, y)$ to return maximum of x and y. Here, the smallest solution expression is `if` $x \le y$ `then` y `else` x.
Basic Enumerative Strategy. We explain the basic enumerative algorithm [24] using Table 1. The enumerative algorithm maintains a set of input points `pts` (initially empty), and proceeds in rounds. In each round the algorithm proposes a solution expression which is correct on all of `pts`. If this expression is correct on all inputs (verified by an SMT solver) the algorithm returns it. Otherwise, a counter-example input point is added to `pts`.

The algorithm generates the candidate solution expression by enumerating expressions generated by the grammar in order of size. In the first round, the candidate expression proposed is the first expression generated (the expression 0) as `pts` is empty. Attempting to verify the correctness of this expression, yields a counter-example point $\{x \mapsto 1, y \mapsto 0\}$. In the second round, the expression 0 is incorrect on the point, and the next expression to be correct on all of `pts` (the expression 1) is proposed. This fails to verify as well, and yields the

Round no.	Enumerated Terms	Enumerated Predicates	Candidate Expression	Point added
1	0	0	$\emptyset$	$\{x \mapsto 1, y \mapsto 0\}$
2	0, 1	1	$\emptyset$	$\{x \mapsto 0, y \mapsto 2\}$
3	0, 1, x, y	$0 \leq 0, \ldots 0 \leq y,$ $1 \leq 0, \ldots 1 \leq y$	if $1 \leq y$ then y else 1	$\{x \mapsto 2, y \mapsto 0\}$
4	0, 1, x, y	$0 \leq 0, \ldots x \leq y$	if $x \leq y$ then y else x	

Table 2: Example run of the divide-and-conquer enumerative algorithm

counter-example point $\{x \mapsto 0, y \mapsto 2\}$. In the third round, all expressions of size 1 are incorrect on at least one point in pts, and the algorithm moves on to enumerate larger expressions. After several rounds, the algorithm eventually generates the expression if $x \leq y$ then y else x which the SMT solver verifies to be correct. Note that to get the correct expression, the algorithm had to generate a large number (in this case, hundreds of expressions). In general, the number of expressions grows exponentially with size. Thus, the enumerative solver fails to scale to problems which only admit solutions representable by large expressions. **Divide and Conquer Enumeration**. In the above example, though the solution is large, the individual components (terms x and y, and predicate $x \leq y$) are rather small and can be quickly enumerated. The divide-and-conquer approach enumerates terms and predicates separately, and attempts to combine them into a conditional expression. We explain this idea using an example (see Table 2).

Similar to the basic algorithm, the divide-and-conquer algorithm maintains a set of points pts, and works in rounds. The first two rounds are similar to the run of the basic algorithm. In contrast to the basic algorithm, the enumeration stops in the third round after 0, 1, x, and y are enumerated – the terms 1 and y are correct on $\{x \mapsto 1, y \mapsto 0\}$ and $\{x \mapsto 0, y \mapsto 2\}$, respectively, and thus together "cover" all of pts. Now, to propose an expression, the algorithm starts enumerating predicates until it finds a sufficient number of predicates to generate a conditional expression using the previously enumerated terms. The terms and predicates are combined into conditional expression by learning decision trees (see Section 4.2). The candidate expression proposed in the third round is if $1 \leq y$ then y else x and the counter-example generated is $\{x \mapsto 2, y \mapsto 0\}$ (see table). Proceeding further, in the fourth round, the correct expression is generated. Note that this approach only generates 4 terms and 11 predicates in contrast to the basic approach which generates hundreds of expressions.

3 Problem Statement and Background

Let us fix the function to be synthesized f and its formal parameters params. We write range(f) to denote the range of f. The term *point* denotes a valuation of params, i.e., a point is an input to f.

Example 1. For the running example in this section, we consider a function to be synthesized f of type $\mathbb{Z} \times \mathbb{Z} \to \mathbb{Z}$ with the formal parameters params $= \{x, y\}$. Points are valuations of x and y. For example, $\{x \mapsto 1, y \mapsto 2\}$ is a point.

Specifications. SMT formulae have become the standard formalism for specifying semantic constraints for synthesis. In this paper, we fix an arbitrary theory $\mathcal{T}$ and denote by $\mathcal{T}[\mathsf{symbols}]$, the set of $\mathcal{T}$ terms over the symbols $\mathsf{symbols}$. A *specification* Φ is a logical formula in a theory $\mathcal{T}$ over standard theory symbols and the function to be synthesized f. An expression e *satisfies* Φ ($e \models \Phi$) if instantiating the function to be synthesized f by e makes Φ valid.

Example 2. Continuing the running example, we define a specification $\Phi \equiv \forall x, y :$ $f(x,y) \geq x \wedge f(x,y) \geq y \wedge f(x,y) = x \vee f(x,y) = y$. The specification states that f maps each pair x and y to a value that is at least as great as each of them and equal to one of them, i.e., the maximum of x and y.

Grammars. An *expression grammar* G is a tuple $\langle \mathcal{N}, S, \mathcal{R} \rangle$ where: (a) the set $\mathcal{N}$ is a set of non-terminal symbols, (b) the non-terminal $S \in \mathcal{N}$ is the initial non-terminal, (c) $\mathcal{R} \subseteq \mathcal{N} \times \mathcal{T}[\mathcal{N} \cup \mathsf{params}]$ is a finite set of rewrite rules that map $\mathcal{N}$ to $\mathcal{T}$-expressions over non-terminals and formal parameters. We say that an expression e *rewrites to* an incomplete expression e' (written as $e \rightarrow_G e'$) if there exists a rule $R = (N, e'') \in \mathcal{R}$ and e' is obtained by replacing one occurrence of N in e by e''. Let $\rightarrow_G^*$ be the transitive closure of $\rightarrow$. We say that an expression $e \in \mathcal{T}[\mathsf{params}]$ is *generated* by the grammar G (written as $e \in [\![G]\!]$) if $S \rightarrow_G^* e$. Note that we implicitly assume that all terms generated by the grammar have the right type, i.e., are of the type $\mathsf{range}(f)$.

Example 3. For the running example, we choose the following grammar. The set of non-terminals is given by $\mathcal{N} = \{S, T, C\}$ with the initial non-terminal being S. The rules of this grammar are $\{(S, T), (S, \texttt{if } C \texttt{ then } S \texttt{ else } S)\} \cup \{(T, x), (T, y), (T, 1), (T, 0), (T, T + T)\} \cup \{(C, T \leq T), (C, C \wedge C), (C, \neg C)\}$. This is the standard linear integer arithmetic grammar used for many SyGuS problems. This grammar is equivalent to the one from Figure 1.

The Syntax-Guided Synthesis Problem. An instance of the SyGuS problem is given by a pair $\langle \Phi, G \rangle$ of specification and grammar. An expression e is a solution to the instance if $e \models \Phi$ and $e \in [\![G]\!]$.

Example 4. Continuing the running example, for the specification Φ from Example 2 and the grammar from Example 3, one of the solution expressions is given by $f(x,y) \equiv \texttt{if } x \leq y \texttt{ then } y \texttt{ else } x$.

From our definitions, it is clear that we restrict ourselves to a version of the SyGuS problem where there is exactly one unknown function to be synthesized, and the grammar does not contain $\texttt{let}$ rules. Further, we assume that our specifications are *point-wise*. Intuitively, if a specification is point-wise, if it just relates the input point to its output, and not the outputs of different inputs.

Here, we use a simple syntactic notion of point-wise specifications for convenience. However, our techniques can be generalized to any notion of point-wise specifications. Formally, we say that a specification is *plainly separable* if can be rewritten into a conjunctive normal form where each clause is either (a) a tautology, or (b) every appearing application of the function to be synthesized f have the same arguments.

Example 5. The specification for our running example $\Phi \equiv f(x,y) \geq x \wedge f(x,y) \geq y \wedge f(x,y) = x \vee f(x,y) = y$ is plainly separable. For example, this implies that the value of $f(1,2)$ can be chosen irrespective of the value of f on any other point. On the other hand, a specification such as $f(x,y) = 1 \implies f(x+1,y) = 1$ is neither plainly separable nor point-wise. The value of $f(1,2)$ cannot be chosen independently of the value of $f(0,2)$.

The above restrictions make the SyGuS problem significantly easier. However, a large fraction of problems do fall into this class. Several previous works address this class of problem (see, for example, [3, 15, 19]).

Simply separable specifications allow us to define the notion of an expression e satisfying a specification Φ on a point pt. Formally, we say that $e \models \Phi \downarrow$ pt if e satisfies the specification obtained by replacing each clause C in Φ by $Pre_C(\text{pt}) \implies C$. Here, the premise $Pre_C(\text{pt})$ is given by $\bigwedge_{p \in \text{params}} Arg_C(p) = \text{pt}[p]$ where Arg_C is the actual argument corresponding to the formal parameter p in the unique invocation of f that occurs in C. We extend this definition to sets of points as follows: $e \models \Phi \downarrow \text{pts} \Leftrightarrow \bigwedge_{\text{pt} \in \text{pts}} e \models \Phi \downarrow \text{pt}$.

Example 6. For the specification Φ of the running example, the function given by $f(x,y) \equiv x + y$ is correct on the point $\{x \mapsto 0, y \mapsto 3\}$ and incorrect on the point $\{x \mapsto 1, y \mapsto 2\}$

3.1 The Enumerative Solver

Algorithm 1 Enumerative Solver

Require: Grammar $G = \langle \mathcal{N}, S, \mathcal{R} \rangle$
Require: Specification Φ
Ensure: e s.t. $e \in [\![G]\!] \wedge e \models \Phi$
1: $\text{pts} \leftarrow \emptyset$
2: **while** true **do**
3:    **for** $e \in \text{ENUMERATE}(G, \text{pts})$ **do**
4:       **if** $e \not\models \Phi \downarrow \text{pts}$ **then continue**
5:       $\text{cexpt} \leftarrow \text{verify}(e, \Phi)$
6:       **if** $\text{cexpt} = \bot$ **then return** e
7:       $\text{pts} \leftarrow \text{pts} \cup \text{cexpt}$

The principal idea behind the enumerative solver is to enumerate all expressions from the given syntax with some pruning. Only expressions that are distinct with respect to a set of concrete input points are enumerated.

The full pseudo-code is given in Algorithm 1. Initially, the set of points is set to be empty at line 1. In each iteration, the algorithm calls the ENUMERATE procedure[1] which returns the next element from a (possibly infinite) list of expressions such that no two expressions in this list evaluate to the same values at every point $\text{pt} \in \text{pts}$ (line 3). Every expression e in this list is then verified, first on the set of points (line 4) and then fully (line 5). If the expression e is correct, it is returned (line 6). Otherwise, we pick a counter-example input point (*i.e.*, an input on which e is incorrect) and add it to the set of points and repeat (line 7). A full description of the ENUMERATE procedure can be found in Appendix B.

Theorem 1. *Given a SyGuS instance (Φ, G) with at least one solution expression, Algorithm 1 terminates and returns the smallest solution expression.*

---
[1] Note that ENUMERATE is a coprocedure. Unfamiliar readers may assume that each call to ENUMERATE returns the next expression from an infinite list of expressions.

Fig. 2: Scalability of Enumeration

Features and Limitations. The enumerative algorithm performs surprisingly well considering its simplicity on small to medium sized benchmarks (see [2, 24]). Further, due to the guarantee of Theorem 1 that the enumerative approach produces small solutions, the algorithm is capable of generalizing from specifications that are input-output examples.

However, enumeration quickly fails to scale with growing size of solutions. Figure 2 shows the time taken (in seconds) to generate all expressions up to a given size for the grammar from Figure 1. As can be seen from the graph, the time taken grows exponentially with the size.

4   The Divide-and-Conquer Enumeration Algorithm

Conditional Expression Grammars. We introduce conditional expressions grammars that separate an expression grammar into two grammars that generate: (a) the return value expression, and (b) the conditionals that decide which return value is chosen. These generated return values (terms) and conditionals (predicates) are combined using if-then-else conditional operators.

A *conditional expression grammar* is a pair of grammars $\langle G_T, G_P \rangle$ where: (a) the *term grammar* G_T is an expression grammar generating terms of type range(f); and (b) the *predicate grammar* G_P is an expression grammar generating boolean terms. The set of expressions $[\![\langle G_T, G_P \rangle]\!]$ generated by $\langle G_T, G_P \rangle$ is the smallest set of expressions $\mathcal{T}[\text{params}]$ such that: (a) $[\![G_T]\!] \subseteq [\![\langle G_T, G_P \rangle]\!]$, and (b) $e_1, e_2 \in [\![\langle G_T, G_P \rangle]\!] \wedge p \in [\![G_P]\!] \implies$ if p then e_1 else $e_2 \in [\![\langle G_T, G_P \rangle]\!]$ Most commonly occurring SyGuS grammars in practice can be rewritten as conditional expression grammars automatically.

Example 7. The grammar from Example 3 is easily decomposed into a conditional expression grammar $\langle G_T, G_P \rangle$ where: (a) the term grammar G_T contains only the non-terminal T, and the rules for rewriting T. (b) the predicate grammar G_P contains the two non-terminals $\{T, C\}$ and the associated rules.

Decision Trees. We use the concept of decision trees from machine learning literature to model conditional expressions. Informally, a decision tree DT maps *samples* to *labels*. Each internal node in a decision tree contains an *attribute* which may either hold or not for each sample, and each leaf node contains a label. In our setting, labels are terms, attributes are predicates, and samples are points.

To compute the label for a given point, we follow a path from the root of the decision tree to a leaf, taking the left (resp. right) child at each internal node if the attribute holds (resp. does not hold) for the sample. The required label is the label at the leaf. We do not formally define decision trees, but instead refer the reader to a standard text-book (see, for example, [4]).

Example 8. Figure 3 contains a decision tree in our setting, i.e., with attributes being predicates and labels being terms. To compute the associated label with the point $\mathsf{pt} \equiv \{x \mapsto 2, y \mapsto 0\}$: (a) we examine the predicate at the root node, i.e., $y \leq 0$ and follow the left child as the predicate hold for pt; (b) examine the predicate at the left child of the root node, i.e, $x \leq y$ and follow the right child as it does not hold; and (c) return the label of the leaf $x + y$.


Fig. 3: Sample decision tree

The expression $\mathsf{expr}(DT)$ corresponding to a decision tree DT is: (a) the label of the root node if the tree is a single leaf node; and (b) if p then $\mathsf{expr}(DT_L)$ else $\mathsf{expr}(DT_Y)$ where p is the attribute of the root node, and DT_L and DT_Y are the left and right children, otherwise.

Decision tree learning is a technique that learns a decision tree from a given set of samples. A decision tree learning procedure is given: (a) a set of samples (points), (b) a set of labels (terms), along with a function that maps a label to the subset of samples which it covers; and (c) a set of attributes (predicates). A sound decision tree learning algorithm returns a decision tree DT that classifies the points correctly, i.e., for every sample pt, the label associated to it by the decision tree covers the point. We use the notation `LearnDT` to denote a generic sound decision tree learning procedure. The exact procedure we use for decision tree learning in Section 4.2.

4.1 Algorithm

Algorithm 2 presents the full divide-and-conquer enumeration algorithm for synthesis. Like Algorithm 1, the divide-and-conquer algorithm maintains a set of points pts, and in each iteration: (a) computes a candidate solution expression e (lines 3-10); (b) verifies and returns e if it is correct (lines 10 and 11); and (c) otherwise, adds the counter-example point into the set pts (line 12).

However, the key differences between Algorithm 2 and Algorithm 1 are in the way the candidate solution expression e is generated. The generation of candidate expressions is separated into two separate steps.

Term solving. Instead of searching for a single candidate expression that is correct on all points in pts, Algorithm 2 maintains a set of candidate terms terms. We say that a term t covers a point $\mathsf{pt} \in \mathsf{pts}$ if $t \models \Phi \downarrow \mathsf{pt}$. The set of points that a term covers is computed and stored in $\mathsf{cover}[t]$ (line 15). Note that the algorithm does not store terms that cover the same set of points as already generated terms (line 16). When the set of terms terms covers all the points in pts, i.e., for each $\mathsf{pt} \in \mathsf{pts}$, there is at least one term that is correct on pt, the term enumeration is stopped (while-loop condition in line 4).

Unification and Decision Tree Learning.. In the next step (lines 6-9), we generate a set of predicates preds that will be used as conditionals to combine the terms from terms into if-then-else expressions. In each iteration, we attempt to learn a decision tree that correctly label each point $\mathsf{pt} \in \mathsf{pts}$ with a term t such that $\mathsf{pt} \in \mathsf{cover}[t]$. If such a decision tree DT exists, the conditional expression $\mathsf{expr}(DT)$ is correct on all points, i.e., $\mathsf{expr}(DT) \models \Phi \downarrow \mathsf{pts}$. If a decision tree does not exist, we generate additional terms and predicates and retry.

Algorithm 2 DCSolve: The divide-and-conquer enumeration algorithm

---

Require: Conditional expression grammar $G = \langle G_T, G_P \rangle$
Require: Specification Φ
Ensure: Expression e s.t. $e \in [\![G]\!] \wedge e \models \Phi$

1: pts ← ∅
2: **while** true **do**
3:     terms ← ∅; preds ← ∅; cover ← ∅; $DT = \bot$
4:     **while** $\bigcup_{t \in \text{terms}}$ cover$[t] \neq$ pts **do**   ▷ Term solver
5:         terms ← terms ∪ NEXTDISTINCTTERM(pts, terms, cover)
6:     **while** $DT = \bot$ **do**   ▷ Unifier
7:         terms ← terms ∪ NEXTDISTINCTTERM(pts, terms, cover)
8:         preds ← preds ∪ ENUMERATE(G_P, pts)
9:         DT ← LEARNDT(terms, preds)
10:     e ← expr(DT); cexpt ← verify(e, Φ)   ▷ Verifier
11:     **if** cexpt $= \bot$ **then return** e
12:     pts ← pts ∪ cexpt
13: **function** NEXTDISTINCTTERM(pts, terms, cover)
14:     **while** *True* **do**
15:         t ← ENUMERATE(G_T, pts); cover$[t]$ ← {pt | pt ∈ pts $\wedge t \models \Phi \downarrow$ pt}
16:         **if** $\forall t' \in$ terms : cover$[t] \neq$ cover$[t']$ **then return** t

---

Remark 1. In line 7, we generate additional terms even though terms is guaranteed to contain terms that cover all points. This is required to achieve semicompleteness, i.e., without this, the algorithm might not find a solution even if one exists. See the appendix for an example.

Theorem 2. *Algorithm 2 is a sound for the SyGuS problem. Further, assuming a sound and complete* LearnDT *procedure, if there exists a solution expression, Algorithm 2 is guaranteed to find it.*

The proof of the above theorem is similar to the proof of soundness and partial-completeness for the original enumerative solver. The only additional assumption is that the LearnDT decision tree learning procedure will return a decision tree if one exists. We present such a procedure in the next section.

4.2   Decision Tree Learning

The standard multi-label decision tree learning algorithm (based on ID3 [18]) is presented in Algorithm 3. The algorithm first checks if there exists a single label (i.e., term) t that applies to all the points (line 1). If so, it returns a decision tree with only one leaf whose label is t (line 1). Otherwise, it picks the best predicate p to split on based on some heuristic (line 3). If no predicates are left, there exists no decision tree, and the algorithm returns $\bot$ (line 2). It the recursively computes the left and right sub-trees for the set of points on which p holds and does not hold, respectively (lines 4 and 5). The final decision tree is returned as a tree with a root (with attribute p), and positive and negative edges to the roots of the left and right sub-trees, respectively.
Information-gain heuristic. The choice of the predicate at line 3 influences the size of the decision tree learned by Algorithm 3, and hence, in our setting, the size of the solution expression generated by Algorithm 2. We use the classical

Algorithm 3 Learning Decision Trees

Require: pts, terms, cover, preds
Ensure: Decision tree DT
 1: **if** $\exists t : \text{pts} \subseteq \text{cover}[t]$ **then return** $LeafNode[\mathcal{L} \leftarrow t]$
 2: **if** $\text{preds} = \emptyset$ **then return** $\bot$
 3: $p \leftarrow$ Pick predicate from preds
 4: $L \leftarrow \text{LEARNDT}(\{\text{pt} \mid p[\text{pt}]\}, \text{terms}, \text{cover}, \text{preds} \setminus \{p\})$
 5: $R \leftarrow \text{LEARNDT}(\{\text{pt} \mid \neg p[\text{pt}]\}, \text{terms}, \text{cover}, \text{preds} \setminus \{p\})$
 6: **return** $InternalNode[\mathcal{A} \leftarrow p, left \leftarrow L, right \leftarrow R]$

---

information gain heuristic to pick the predicates. Informally, the information gain heuristic treats the label as a random variable, and picks to split on the attribute whose value will reveal the most information about the label. We do not describe all aspects of computing information gain, but refer the reader to any standard textbook on machine learning [4]. Given a set of points $\text{pts}' \subseteq \text{pts}$ the entropy $H(\text{pts}')$ is defined in terms of the probability $\mathbb{P}_{\text{pts}'}(label(\text{pt}) = t)$ of a point $\text{pt} \in \text{pt}'$ being labeled with the term t as

$$H(\text{pts}') = -\sum_t \mathbb{P}_{\text{pts}'}(\text{label}(\text{pt}) = t) \cdot \log_2 \mathbb{P}_{\text{pts}'}(\text{label}(\text{pt}) = t)$$

Further, given a predicate $p \in \text{preds}$, the information gain of p is defined as

$$G(p) = \frac{|\text{pts}_Y|}{|\text{pts}|} \cdot H(\text{pts}_Y) + \frac{|\text{pts}_N|}{|\text{pts}|} \cdot H(\text{pts}_N)$$

where $\text{pts}_Y = \{\text{pt} \in \text{pts} \mid p[\text{pt}]\}$ and $\text{pts}_N = \{\text{pt} \in \text{pts} \mid \neg p[\text{pt}]\}$. Hence, at line 3, we compute the value $G(p)$ for each predicate in preds, and pick the one which maximizes $G(p)$.

We use conditional probabilities $\mathbb{P}_{\text{pts}'}(\text{label}(\text{pt}) = t \mid \text{pt})$ to compute the probability $\mathbb{P}_{\text{pts}'}(\text{label}(\text{pt}) = t)$. The assumption we make about the prior distribution is that the likelihood of a given point pt being labelled by a given term t is proportional to the number of points in $\text{cover}[t]$. Formally, we define:

$$\mathbb{P}_{\text{pts}'}(\text{label}(\text{pt}) = t \mid \text{pt}) = \begin{cases} 0 & \text{if } \text{pt} \notin \text{cover}[t] \\ \dfrac{|\text{cover}[t] \cap \text{pts}'|}{\displaystyle\sum_{t' \mid \text{pt} \in \text{cover}[t']} |\text{cover}[t'] \cap \text{pts}'|} & \text{if } \text{pt} \in \text{cover}[t] \end{cases}$$

Now, the unconditional probability of an arbitrary point being labelled with t is given by $\mathbb{P}_{\text{pts}'}(\text{label}(\text{pt}) = t) = \sum_{\text{pt}} \mathbb{P}_{\text{pts}'}(\text{label}(\text{pt}) = t \mid \text{pt}) \cdot \mathbb{P}_{\text{pts}'}(\text{pt})$. Assuming a uniform distribution for picking points, we have that

$$\mathbb{P}_{\text{pts}'}(\text{label}(\text{pt}) = t) = \frac{1}{|\text{pts}|} \cdot \sum_{\text{pt}} \mathbb{P}_{\text{pts}'}(\text{label}(\text{pt}) = t \mid \text{pt})$$

4.3 Extensions and Optimizations

The Anytime Extension.. Algorithm 2 stops enumeration of terms and predicates as soon as it finds a single solution to the synthesis problem. However,

there are cases where due to the lack of sufficiently good predicates, the decision tree and the resulting solution can be large (see Example 9). Instead, we can let the algorithm continue by generating more terms and predicates. This could lead to different, potentially smaller decision trees and solutions.

Example 9. Consider the specification $(x \geq 0 \wedge y \geq 0) \implies (f(x, y) = 1 \Leftrightarrow x + y \leq 2)$ and a run of Algorithm 2 where the terms 0 and 1 are generated. These terms fully cover any set of points for this specification.

Over a sequence of iterations the predicates are generated in order of size. Now, for predicate size 3, the predicates generated include $x = 0$, $x = 1$, $x = 2$, $y \leq 2$, $y \leq 1$, and $y \leq 0$. With these predicates, the decision tree depicted in Figure 4a is learned, and the corresponding conditional expression is correct for the specification. However, if the procedure continues to run after the first solution is generated, predicates of size 4 are generated. Among these predicates, the predicate $x + y \leq 2$ is also generated. With this additional predicate, the decision tree in Figure 4b is generated, leading to the compact solution $f(x, y) \equiv \text{if } x + y \leq 2 \text{ then } 1 \text{ else } 0$.

(a) Decision tree for predicates of size 3

(b) Decision tree for predicates of size 4

Decision Tree Repair. In Algorithm 2, we discard the terms that cover the same set of points as already generated terms in line 16. However, these discarded terms may lead to better solutions than the already generated ones.

Example 10. Consider a run of the algorithm for the running example, where the set pts contains the points $\{x \mapsto 1, y \mapsto 0\}$ and $\{x \mapsto -1, y \mapsto 0\}$. Suppose the algorithm first generates the terms 0 and 1. These terms are each correct on one of the points and are added to terms. Next, the algorithm generates the terms x and y. However, these are not added to terms as x (resp. y) is correct on exactly the same set of points as 1 (resp. 0).

Suppose the algorithm also generates the predicate $x \leq y$ and learns the decision tree corresponding to the expression $e \equiv \text{if } x \leq y \text{ then } 0 \text{ else } 1$. Now, verifying this expression produces a counter-example point, say $\{x \mapsto 1, y \mapsto 2\}$. While the term 0, and correspondingly, the expression e is incorrect on this point, the term y which was discarded as an equivalent term to 0, is correct.

Hence, for a practical implementation of the algorithm we do not discard these terms and predicates, but store them separately in a map $Eq : \text{terms} \rightarrow [\![G_T]\!]$ that maps the terms in terms to an additional set of equivalent terms. At lines 16, if the check for distinctness fails, we instead add the term t to the Eq map. Now, when the decision tree learning algorithm returns an expression that fails to verify and returns a counter-example, we attempt to replace terms and predicates in the decision tree with equivalent ones from the Eq map to make the decision tree behave correctly on the counter-example.

11

Example 11. Revisiting Example 10, instead of discarding the terms x and y, we store them into the *Eq* array, i.e., we set $Eq(0) = \{y\}$ and $Eq(1) = \{x\}$. Now, when the verification of the expression fails, with the counter-example point $\{x \mapsto 1, y \mapsto 2\}$, we check the term that is returned for the counter-example point–here, 0. Now, we whether any term in $Eq(0)$ is correct on the counter-example point–here, the term y. If so, we replace the original term with the equivalent term that is additionally correct on the counter-example point and proceed with verification. Replacing 0 with y in the expression gives us if $x \leq y$ then y else 1. Another round of verification and decision tree repair will lead to replacing the term 1 with x, giving us the final correct solution.

Branch-wise verification. In Algorithm 2, and in most synthesis techniques, an incorrect candidate solution is used to generate 1 counter-example point. However, in the case of conditional expressions and point-wise specifications, each branch (i.e., leaf of the decision tree) can be verified separately. Verifying each branch involves rewriting the specification as in the point-wise verification defined in Section 3 – but instead of adding a premise to each clause asserting that the arguments to the function are equal to a point, we add a premise that asserts that the arguments satisfy all predicates along the path to the leaf. This gives us two separate advantages:
  - We are able to generate multiple counter-examples from a single incorrect expression. This reduces the total number of iterations required, as well as the number of calls to the expensive decision tree learning algorithm.
  - It reduces the complexity of each call to the verifier in terms of the size of the SMT formula to be checked. As verification procedures generally scale exponentially with respect to the size of the SMT formula, multiple simpler verification calls are often faster than one more complex call.
This optimization works very well along with the decision tree repair described above as we can verify and repair each branch of the decision tree separately.

Example 12. Consider the verification of the expression if $x \leq y$ then 0 else 1 for the running example. Instead of running the full expression through the verifier to obtaini one counter-example point, we can verify the branches separately by checking the satisfiability of the formulae $x \leq y \wedge f(x,y) = 0 \wedge \neg\Phi$ and $\neg(x \leq y) \wedge f(x,y) = 1 \wedge \neg\Phi$. This gives us two separate counter-example points.

5 Evaluation

We built a prototype SyGuS solver that uses the divide-and-conquer enumerative algorithm. The tool consists of 3000 lines of Python code implementing the enumeration, and 3000 lines of C++ code implementing the decision tree learning. All experiments were executed on a 32-core machine with four Intel Xeon E7-4820 processors and 128 GB of memory, with a one hour time limit.
Goals. We seek to empirically evaluate how our synthesis algorithm compares to other state-of-the-art synthesis techniques along the following dimensions: (a) *Performance*: How quickly can the algorithms arrive at a correct solution?

(b) *Quality*: How *good* are the solutions produced by the algorithms? We use compactness of solutions as a metric for the quality of solutions. (c) *Effect of continued execution*: How significant is the improvement in the quality of the solutions generated given an additional (but fixed) time budget.

(a) Total number of ICFP benchmarks solved by DCSolve as a function of elapsed time. Interpretation: The y-coordinate of every point on the black (resp. red) plot gives the number of benchmarks for which the DCSolve algorithm can produce the first (resp. smallest) solution within the time indicated by the x-coordinate of the point. All times reported here are *per benchmark* times.
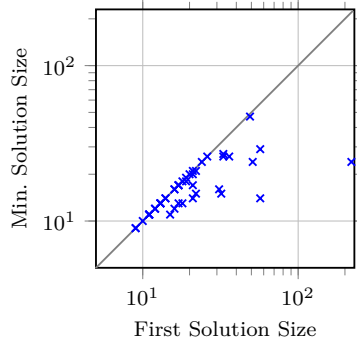
(b) Scatter plot of first vs. minimum size solutions. Interpretation: Each point represents a benchmark. The x-coordinate of each point gives the size of the *first* solution produced by DCSolve. The y-coordinate gives the size of the *smallest* solution produced by DCSolve within 3600 seconds. Points below the $x = y$ line correspond to benchmarks which benefit from continued execution, beyond the first solution.

Fig. 5: Summary graphs of experimental results on the ICFP benchmarks

Benchmarks. We considered two classes of SyGuS benchmarks taken from the SyGuS competition 2015. The first class that we considered consists of the ICFP benchmarks. These are bit-twiddling problems with (a) very specific syntactic restrictions on the operations allowed, and (b) behaviour specified using only concrete input-output examples To the best of our knowledge, no existing SyGuS solver has been able to solve these benchmarks. Detailed experimental results on the ICFP benchmarks can be found in Tables 3 and 4 in Appendix A.

The second class of benchmarks that we considered compute the maximum of n integers, where n is a parameter. Detailed results on these benchmarks can be found in Table 5 in Appendix A.

The State-of-the-art. Empirically, black-box algorithms are able to generalize well from underconstrained correctness specification, whereas it is not clear how white-box solvers can generalize from such specifications, given their heavy reliance on the specifications being complete. We note that ESOLVER, which employs an optimized version of the basic enumerative algorithm discussed earlier won the 2014 SyGuS competition and came in second in the 2015 SyGuS competition, losing to the CVC4 solver, which came in first.

5.1 Discussion

Performance. As Tables 3 and 4, as well as the black plot in Figure 5a demonstrate, our algorithm was able to produce a solution for 47 out of the 50 ICFP benchmarks, within an hour. We emphasize that we know of no other solver that has been able to provide satisfactory solutions to

these benchmarks. The solutions produced by our divide and conquer algorithm, while being compact, are often of size larger than 20, indicating that the solutions are beyond the reach of a basic enumerative strategy. Of the 47 solved instances, only 5 benchmarks take longer than 10 minutes, as seen in Figure 5a.

Table 5 indicates that our algorithm is competitive with the state-of-the-art CVC4 solver on the parameterized "max of n numbers" benchmarks, for smaller values of n. For larger values of n, our algorithm is not as performant. Our investigations indicated that in the performance drops are largely due to the large number of concrete counterexample points that the algorithm maintains, which in turn leads to the decision tree learning algorithms taking longer.

Quality of Solutions. For ICFP benchmarks, Tables 3 and 4 indicate that most of the solutions that our algorithm generates are rather compact, and are of a size that are comprehensible to a human. We do not know if solutions smaller than these exist. The CVC4 solver has been able to synthesize solutions for some of the ICFP benchmarks, *but only when syntactic restrictions are lifted*. However, the solutions generated are large case-splits corresponding to the input-output examples that form the constraints for the problem. This is because CVC4 is based on model-based quantifier instantiation, and each input-output example is an instantiation of the quantifier. The solutions are extremely large, and are not the intended or desired solutions. Our solutions are compact expressions that perform some computation on the input, to return an output.

In the case of the "max" benchmarks, we again see from Table 5 that our algorithm performs well for small parameter values, often beating the solutions obtained from the CVC4 solver in terms of size and number of comparisons. At larger parameter values, the solutions are still reasonably compact and not excessively large. The reason for the somewhat larger solutions for larger parameter values is that the decision tree learning algorithm is not optimal, *i.e.*, it does not guarantee the most compact decision tree. Instead a greedy heuristic is used, which rapidly yields suboptimal solutions as the number of samples (counterexample points) increases. We intend to explore how we can adapt other heuristics [15], into the setting of our problem, in future work.

Effect of Continued Execution. For about half the ICFP benchmarks, we were able to obtain a more compact solution by letting the algorithm continue execution after the first solution was discovered (last column in Tables 3 and 4, and Figure 5b). Further, the difference in the first and smallest solutions is sometimes very significant. For example, in the case of the "icfp_95_100" benchmark, we see the size of the solution go down from 220 to 24. An interesting phenomenon that we observed was that while the size of the decision tree almost always went down with time, the size of the solutions sometimes increased. This is because the algorithm generated larger terms and predicates over time, increasing the size of the labels and attributes in each node of the decision tree.

Overall, our experimental evaluation suggests that: (a) The DCSolve algorithm is able to quickly learn compact solutions, and generalizes well from input-output examples. (b) The anytime nature of the DCSolve algorithm is often useful in reducing the size of the computed solution; (c) The DCSolve algorithm works competently on problems from different classes of benchmarks.

14

6   Concluding Remarks

Related Work. Program synthesis has seen a revived interest in the last decade, starting from the SKETCH framework [22, 23] which proposed the counterexample guided inductive synthesis (CEGIS). Almost all synthesis algorithms proposed in recent literature can be viewed as an instantiation of CEGIS. Synthesis of string manipulating programs using input-output examples has found applications in Microsoft's FlashFill [8], and the ideas have been generalized for other domains in a meta-synthesis framework called FlashMeta [17]. Other recent work in the area of program synthesis have used type-theoretic approaches [6, 10, 16] for program completion and for generating code snippets. Synthesis of recursive programs and data structure manipulating code has also been studied extensively in recent years [1, 5, 13, 14]. Lastly, synthesis techniques based on decision trees have been used to learn program invariants [7].

In the area of SyGuS, solvers based on enumerative search [24], stochastic search [2, 21] and symbolic search [9, 12] were among the first solvers developed. The SKETCH approach has also been used to develop SyGuS solvers [11]. Alchemist [20] is another solver that is quite competitive on benchmarks in the linear arithmetic domains. More recently, white box solvers like the CVC4 solver [19] and the unification based solver [3] have also been developed.

The enumerative synthesis algorithm used by ESOLVER [2, 24] and the work on using decision trees for piece-wise functions [15] are perhaps the most closely related to the work described in this paper. We have already discussed at length the shortcomings of ESOLVER that our algorithm overcomes. The approach for learning piece-wise functions [15] also uses decision trees. While the presented framework is generic, the authors instantiate and evaluate it only for the linear arithmetic domain with a specific grammar. In DCSolve, neither the decision tree learning algorithm, nor the enumeration is domain-specific, making DCSolve a domain and grammar agnostic algorithm. The algorithm presented in [15] can easily learn large constants in the linear integer domain. This is something that enumerative approaches, including DCSolve struggle to do. The heuristics used for decision tree learning are different; in [15], the authors use a heuristic based on hitting sets, while we use a information gain heuristic with cover-based priors.

Conclusion. This paper has presented a new enumerative algorithm to solve instances of the Syntax-Guided Synthesis (SyGuS) problem. The algorithm overcomes the shortcomings of a basic enumerative algorithm by using enumeration to only learn small expressions which are correct on subsets of the inputs. These expressions are then used to form a conditional expression using Boolean combinations of enumerated predicates using decision trees. We have demonstrated the ability of the algorithm to generalize from input-output examples by solving most of the previously unsolved ICFP benchmarks, which are part of the SyGuS benchmark suite. The algorithm is generic, efficient, produces compact solutions, and is *anytime* — in that continued execution can potentially produce more compact solutions.

References

[1] Aws Albarghouthi, Sumit Gulwani, and Zachary Kincaid. Recursive Program Synthesis. In *Computer Aided Verification - 25th International Conference, CAV 2013, Saint Petersburg, Russia, July 13–19, 2013. Proceedings*, pages 934–950, 2013.

[2] Rajeev Alur, Rastislav Bodík, Garvit Juniwal, Milo M. K. Martin, Mukund Raghothaman, Sanjit A. Seshia, Rishabh Singh, Armando Solar-Lezama, Emina Torlak, and Abhishek Udupa. Syntax-guided Synthesis. In *Formal Methods in Computer-Aided Design, FMCAD 2013, Portland, OR, USA, October 20–23, 2013*, pages 1–8, 2013.

[3] Rajeev Alur, Pavol Cerný, and Arjun Radhakrishna. Synthesis Through Unification. In *Computer Aided Verification - 27th International Conference, CAV 2015, San Francisco, CA, USA, July 18–24, 2015, Proceedings, Part II*, pages 163–179, 2015.

[4] Christopher M. Bishop. *Pattern Recognition and Machine Learning (Information Science and Statistics)*. Springer-Verlag New York, Inc., Secaucus, NJ, USA, 2006. ISBN 0387310738.

[5] John K. Feser, Swarat Chaudhuri, and Isil Dillig. Synthesizing Data Structure Transformations from Input-output Examples. In *Proceedings of the 36th ACM SIGPLAN Conference on Programming Language Design and Implementation, Portland, OR, USA, June 15-17, 2015*, pages 229–239, 2015.

[6] Jonathan Frankle, Peter-Michael Osera, David Walker, and Steve Zdancewic. Example-directed Synthesis: A Type-theoretic Interpretation. In *Proceedings of the 43rd Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL 2016, St. Petersburg, FL, USA, January 20 – 22, 2016*, pages 802–815, 2016.

[7] Pranav Garg, Daniel Neider, P. Madhusudan, and Dan Roth. Learning Invariants using Decision Trees and Implication Counterexamples. In *Proceedings of the 43rd Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL 2016, St. Petersburg, FL, USA, January 20 - 22, 2016*, pages 499–512, 2016.

[8] Sumit Gulwani. Automating String Processing in Spreadsheets using Input-output Examples. In *Proceedings of the 38th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL 2011, Austin, TX, USA, January 26–28, 2011*, pages 317–330, 2011.

[9] Sumit Gulwani, Susmit Jha, Ashish Tiwari, and Ramarathnam Venkatesan. Synthesis of Loop-free Programs. In *Proceedings of the 32nd ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI 2011, San Jose, CA, USA, June 4–8, 2011*, pages 62–73, 2011.

[10] Tihomir Gvero, Viktor Kuncak, Ivan Kuraj, and Ruzica Piskac. Complete Completion using Types and Weights. In *ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI '13, Seattle, WA, USA, June 16-19, 2013*, pages 27–38, 2013.

[11] Jinseong Jeon, Xiaokang Qiu, Armando Solar-Lezama, and Jeffrey S. Foster. Adaptive Concretization for Parallel Program Synthesis. In *Computer Aided Verification - 27th International Conference, CAV 2015, San Francisco, CA, USA, July 18-24, 2015, Proceedings, Part II*, pages 377–394, 2015.

[12] Susmit Jha, Sumit Gulwani, Sanjit A. Seshia, and Ashish Tiwari. Oracle-guided Component-based Program Synthesis. In *Proceedings of the 32nd ACM/IEEE*

International Conference on Software Engineering - Volume 1, ICSE 2010, Cape Town, South Africa, 1–8 May 2010, pages 215–224, 2010.

[13] Etienne Kneuss, Ivan Kuraj, Viktor Kuncak, and Philippe Suter. Synthesis Modulo Recursive Functions. In *Proceedings of the 2013 ACM SIGPLAN International Conference on Object Oriented Programming Systems Languages & Applications, OOPSLA 2013, part of SPLASH 2013, Indianapolis, IN, USA, October 26-31, 2013*, pages 407–426, 2013.

[14] Viktor Kuncak, Mikaël Mayer, Ruzica Piskac, and Philippe Suter. Comfusy: A Tool for Complete Functional Synthesis. In *Computer Aided Verification, 22nd International Conference, CAV 2010, Edinburgh, UK, July 15-19, 2010. Proceedings*, pages 430–433, 2010.

[15] Parthasarathy Madhusudan, Daniel Neider, and Shambwaditya Saha. Synthesizing Piece-wise Functions by Learning Classifiers. In *Tools and Algorithms for the Construction and Analysis of Systems - 21^{st} International Conference, TACAS 2016, Held as Part of the European Joint Conferences on Theory and Practice of Software, ETAPS 2016, Eindhoven, Netherlands, April 2 – 8, 2016. Proceedings*, 2016.

[16] Peter-Michael Osera and Steve Zdancewic. Type-and-example-directed Program Synthesis. In *Proceedings of the 36^{th} ACM SIGPLAN Conference on Programming Language Design and Implementation, Portland, OR, USA, June 15–17, 2015*, pages 619–630, 2015.

[17] Oleksandr Polozov and Sumit Gulwani. FlashMeta: A Framework for Inductive Program Synthesis. In *Proceedings of the 2015 ACM SIGPLAN International Conference on Object-Oriented Programming, Systems, Languages, and Applications, OOPSLA 2015, part of SLASH 2015, Pittsburgh, PA, USA, October 25–30, 2015*, pages 107–126, 2015.

[18] J. Ross Quinlan. Induction of Decision Trees. *Machine Learning*, 1(1):81–106, 1986.

[19] Andrew Reynolds, Morgan Deters, Viktor Kuncak, Cesare Tinelli, and Clark W. Barrett. Counterexample-Guided Quantifier Instantiation for Synthesis in SMT. In *Computer Aided Verification - 27^{th} International Conference, CAV 2015, San Francisco, CA, USA, July 18–24, 2015, Proceedings, Part II*, pages 198–216, 2015.

[20] Shambwaditya Saha, Pranav Garg, and P. Madhusudan. Alchemist: Learning Guarded Affine Functions. In *Computer Aided Verification - 27^{th} International Conference, CAV 2015, San Francisco, CA, USA, July 18-24, 2015, Proceedings, Part I*, pages 440–446, 2015.

[21] Eric Schkufza, Rahul Sharma, and Alex Aiken. Stochastic Superoptimization. In *Architectural Support for Programming Languages and Operating Systems, ASPLOS 2013, Houston, TX, USA - March 16 – 20, 2013*, pages 305–316, 2013.

[22] Armando Solar-Lezama, Liviu Tancau, Rastislav Bodík, Sanjit A. Seshia, and Vijay A. Saraswat. Combinatorial Sketching for Finite Programs. In *Proceedings of the 12th International Conference on Architectural Support for Programming Languages and Operating Systems, ASPLOS 2006, San Jose, CA, USA, October 21-25, 2006*, pages 404–415.

[23] Armando Solar-Lezama, Rodric M. Rabbah, Rastislav Bodík, and Kemal Ebcioğlu. Programming by Sketching for Bit-streaming Programs. In *Proceedings of the ACM SIGPLAN 2005 Conference on Programming Language Design and Implementation, Chicago, IL, USA, June 12-15, 2005*, pages 281–294, 2005.

[24] Abhishek Udupa, Arun Raghavan, Jyotirmoy V. Deshmukh, Sela Mador-Haim, Milo M. K. Martin, and Rajeev Alur. TRANSIT: Specifying Protocols with Concolic Snippets. In *ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI 2013, Seattle, WA, USA, June 16–19, 2013*, pages 287–296, 2013.

Appendix A   Detailed Experimental Results

Benchmark	Size of first	Time to first	DT Size for first	Size of min	Time to min	DT Size for min	Min first?
icfp_103_10	33	12.1	7	26	427.4	5	×
icfp_104_10	17	1.1	5	17	30.9	5	×
icfp_105_1000	15	35.9	5	11	134.9	3	×
icfp_105_100	16	3.1	5	12	9.1	3	×
icfp_113_1000	11	53.6	3	11	53.6	3	✓
icfp_114_100	20	241.6	5	21	793.1	5	×
icfp_118_100	31	18.1	7	16	135.8	3	×
icfp_118_10	32	4.1	7	15	9.1	3	×
icfp_125_10	57	8.8	13	14	437.3	3	×
icfp_134_1000	33	1441.1	7	27	2437.4	5	×
icfp_135_100	13	54.0	3	13	54.0	3	✓
icfp_139_10	10	1.3	3	10	1.3	3	✓
icfp_14_1000	–	TO	–	–	TO	–	–
icfp_143_1000	16	1412.5	3	16	1412.5	3	✓
icfp_144_1000	24	1376.8	5	24	1376.8	5	✓
icfp_144_100	36	763.5	7	26	2810.3	5	×
icfp_147_1000	12	1114.3	3	12	1114.3	3	✓
icfp_150_10	51	1.7	11	24	9.3	5	×
icfp_21_1000	21	512.6	5	20	1099.5	5	×
icfp_25_1000	22	151.6	5	15	1571.4	3	×
icfp_28_10	2	0.04	0	2	0.04	0	✓
icfp_30_10	14	14.4	3	14	14.4	3	✓
icfp_32_10	14	9.1	3	14	9.1	3	✓
icfp_38_10	21	5.9	5	14	54.1	3	×
icfp_39_100	12	15.8	3	12	15.8	3	✓

Table 3: Results on the ICFP benchmark suite. All times are in seconds.

Detailed results of running our solver on the ICFP benchmarks are presented in in Tables 3 and 4. For each benchmark, we report (a) the size of the first solution discovered, the time at which it was discovered, and the associated decision tree size; (b) the same details for the minimal solution discovered; and (c) whether the first solution discovered was itself minimal sized solution.

Table 5 demonstrates how our algorithm stacks up against the state-of-the-art CVC4 solver, when evaluated with the "max" benchmarks, which are described in Section 5. The columns indicate the sizes of solutions generated, the time to arrive at the solution, as well as the number of comparisons in the solution obtained by each of the solvers, for different values of the parameter n, indicated in the first column.

Appendix B   Description of the enumerate Procedure

Algorithm 4 describes the ENUMERATE algorithm referred to in Section 3.1. We continue to use the same notation here that was introduced in Section 3.1.

Benchmark	Size of first	Time to first	DT Size for first	Size of min	Time to min	DT Size for min	Min first?
icfp_45_1000	9	15.6	3	9	15.6	3	✓
icfp_45_10	9	0.3	3	9	0.3	3	✓
icfp_5_1000	19	37.4	5	18	2534.9	5	×
icfp_51_10	11	1.9	3	11	1.9	3	✓
icfp_54_1000	11	38.1	3	11	38.1	3	✓
icfp_56_1000	–	TO	–	–	TO	–	–
icfp_64_10	21	21.4	5	21	21.4	5	✓
icfp_68_1000	26	40.1	7	26	40.1	7	✓
icfp_69_10	11	1.1	3	11	1.1	3	✓
icfp_7_1000	17	95.3	5	17	95.3	5	✓
icfp_7_10	16	0.7	5	16	0.7	5	✓
icfp_72_10	13	18.9	3	13	18.9	3	✓
icfp_73_10	18	0.52	5	13	39.6	3	×
icfp_81_1000	21	526.3	5	17	3008.6	3	×
icfp_82_100	13	9.7	3	13	9.7	3	✓
icfp_82_10	13	3.9	3	13	3.9	3	✓
icfp_87_10	19	5.2	5	19	5.2	5	✓
icfp_9_1000	–	TO	–	–	TO	–	–
icfp_93_1000	22	182.1	5	21	318.9	5	×
icfp_94_1000	17	46.8	5	17	46.8	5	✓
icfp_94_100	17	2.4	5	13	498.7	3	×
icfp_95_100	220	270.8	55	24	330.2	7	×
icfp_96_1000	49	2162.3	11	47	2984.1	11	×
icfp_96_10	57	12.4	13	29	12.6	7	×
icfp_99_100	18	321.9	5	18	321.9	5	✓

Table 4: Results on the ICFP benchmark suite (continued from Table 3) . All times are in seconds.

# Args	D&Q Enumeration			CVC4		
	Size	Time (s)	# Comp.	Size	Time (s)	# Comp
2	6	< 0.1	1	6	< 0.1	1
3	16	0.16	3	25	< 0.1	4
4	36	0.77	7	55	< 0.1	9
5	76	6.2	15	96	< 0.1	16
6	156	63.2	31	148	< 0.1	25
7	316	546	63	211	< 0.1	36
8	–	TO	–	285	< 0.1	49
9	–	TO	–	370	0.1	64
10	–	TO	–	466	0.15	81

Table 5: Results on the parameterized "max" benchmarks.

The ENUMERATE procedure maintains a map $\mathsf{prodn} : \mathcal{N} \to \mathbf{2}^{\mathcal{T}[\mathsf{params}]}$ from non-terminals to expressions they can produce. The invariant maintained by the procedure is that every pair of expressions in $\mathsf{prodn}[N]$ is distinct on pts.

Algorithm 4 Enumerating distinct expressions from a grammar

Require: Grammar $G = \langle \mathcal{N}, S, \mathcal{R} \rangle$ and a set of points pts
Ensure: Expressions $\langle e_1, e_2, \ldots \rangle$ s.t. $\forall i < j : |e_i| \leq |e_j| \wedge \exists \mathsf{pt} \in \mathsf{pts} : e_i[\mathsf{pt}] \neq e_j[\mathsf{pt}]$
 1: **for all** $N \in \mathcal{N}$ **do** $\mathsf{prodn}[\mathcal{N}] \leftarrow \emptyset$
 2: **for all** $(N, e) \in \mathcal{R}$ **do**
 3:     **if** $e \in \mathcal{T}[\mathsf{params}]$ **then** $\mathsf{prodn}[N] \leftarrow \mathsf{prodn}[N] \cup \{e\}$
 4: $K \leftarrow 1$
 5: **while** *True* **do**
 6:     **for all** $(N, e) \in \mathcal{R}$ **do**
 7:         $(N_1, \ldots, N_n) \leftarrow$ List of non-terminals occurring in e
 8:         **for all** $(e_1, \ldots, e_n) \in \mathsf{prodn}[N_1] \times \cdots \times \mathsf{prodn}[N_n]$ **do**
 9:             $e^* \leftarrow e[e_1/N_1, \ldots, N_n/e_n]$
10:             **if** $|e^*| \leq K \wedge \forall e' \in \mathsf{prodn}[N].\exists \mathsf{pt} \in \mathsf{pts} : e'[\mathsf{pt}] \neq e^*[\mathsf{pt}]$ **then**
11:                 $\mathsf{prodn}[N] \leftarrow \mathsf{prodn}[N] \cup e^*$
12:                 **if** $N = S$ **then yield** e^*
13:     $K \leftarrow K + 1$

```
S ::= T | if (C) then T else T
T ::= 0 | 1 | 2 | x
C ::= T < 0
```

Fig. 6: Grammar for Appendix C

The algorithm starts by first accumulating into $\mathsf{prodn}[N]$ the expressions that can be produced from N in one step (lines 2-3). Then, for each possible expression size K, it attempts to instantiate each production rule in the grammar with expressions already generated and stored in prodn, to generate new expressions of size at most K. These newly generated expression are checked for distinctness from already generated ones, and if so, added to $\mathsf{prodn}[N]$. The algorithm returns all the expressions produced from the starting non-terminal S.

Appendix C   Algorithm 2 without line 7

Consider the conditional expression grammar given in Figure 6 and the specification $f(x) = x$. Note that the grammar can generate terms equivalent to $ax + b$ where $a, b \geq 0$ and predicates $ax + b < 0$. The solution expression is obviously $f(x) \equiv x$.

Now, in the first round, the algorithm enumerates the first term, i.e., 0. Say counter-example point returned is $\mathsf{pt}_1 = \{x \mapsto 1\}$. In the second round, the algorithm proposes the term 1, and the say counter-example point returned is $\mathsf{pt}_2 = \{x \mapsto 2\}$. In the third round, the algorithm enumerates the terms 1 and 2, which together cover all the previous counter-examples.

Without line 7, any attempts to enumerate predicates and learn decision trees will fail here as no predicate that can be generated by the grammar can distinguish between the pt_1 and pt_2. However, due to line 7, the algorithm will continue to generate more terms while also generating predicates. This will lead to the generation of the term x, and the algorithm will succeed.